

Top Questions Answered: Getting Started with Antigravity

来源: <https://medium.com/google-cloud/top-questions-answered-getting-started-with-antigravity-1b9dd24011d9>

Sitemap

Mastodon## Google Cloud - Community



A collection of technical articles and blogs published or curated by Google Cloud Developer Advocates. The views expressed are those of the authors and don't necessarily reflect those of Google.



During our recent Google Cloud Live episode, “[Getting started with Antigravity](#),” hosts Stephanie Wong, Kevin Hou, and Andy Zhang showcased how Antigravity is lifting the weight of infrastructure management so developers can focus entirely on building. As the team highlighted, Antigravity acts as an agentic colleague — collaborating with you to

plan, research, and execute tasks at an incredibly fast pace.

Our viewers from YouTube, LinkedIn, and dev.to chimed in with some fantastic questions about how Antigravity works under the hood and what it means for the future of coding. Here is a recap of the top questions asked during the stream!

1. How does the Antigravity browser integration work under the hood?

Question from Ryan Swift (dev.to): *I find the Antigravity browser integration to be leagues better than some of the other similar tools out there. What's going on under the hood to actually power this? Is it a special model/training?*

Answer: The seamless browser integration is actually powered by an orchestration of multiple agents. Different types of models — including a specially trained computer-use model — work together to actuate on top of a browser. The specialized model assesses the screen, uses its capabilities to navigate or click, and then reports back to the main model, merging multimodal features seamlessly.

2. Will the coding editor evolve to look less “traditional”?

Question from heckno (dev.to): *Will the editor evolve to look less like a “traditional” editor as AI-assisted editors be more of the norm?*

Answer: Absolutely. Antigravity takes an extreme stance on this with its “Agent Manager,” which the team doesn’t even classify as a traditional editor because you cannot make text edits directly inside it. Instead, the industry is moving towards a future where developers spend less time typing individual lines of code and more time supervising and orchestrating powerful agents that can generate hundreds of lines of code a second.

3. What is the difference between Google AI Studio and Antigravity?

Question from @magicstyleatomic1582 (YouTube): *Why is Google AI Studio different than Antigravity, but the models are the same on the backend?*

Answer: While both leverage the powerful Gemini models under the hood, they serve different stages of the development cycle. AI Studio is an amazing web environment for rapidly building a first iteration of an idea. However, when you need to get deeper into the code, access your local file system, or execute terminal commands, you transition the project onto your local machine using Antigravity.

4. Can you easily roll back changes if an agent makes a mistake?

Question from Adriana Fruchter (LinkedIn): *Do you support versioning or rollback, so we can go back if a generated change doesn't work as expected?*

Answer: Yes! Antigravity has a built-in IDE that supports the standard Git version control you are already used to. Additionally, Antigravity features a “revert” button on each message within the agent conversation. The agent captures the specific file diffs generated between messages, allowing you to easily undo changes up to a specific point and course-correct without losing your broader progress.

5. How easy is it to migrate existing projects?

Question from @saytoga (YouTube): *How can you import projects from VS Code to Antigravity?*

Answer: It is incredibly straightforward. You simply click “open workspace” in Antigravity and select the exact same folder you would normally open in VS Code, and your project will appear instantly.

6. Can I use models other than Gemini?

Question from @kofigame (YouTube): *What other models can be added to Antigravity? Kimi k2, gpt, etc?*

Answer: Antigravity supports an open ecosystem. While the platform is heavily optimized to provide the best possible experience with Google's Gemini models (like Gemini 3.0), it also supports models from OpenAI and Anthropic.

7. How does Antigravity handle security and access control?

Question from Adriana Fruchter (LinkedIn): *As we move into agent-based workflows, is Antigravity thinking about how to handle specifically permissions, access control, and boundaries for what each agent is allowed to do?*

Answer: Safety and security are top priorities (P0) for the Antigravity team. The platform includes several strict guardrails:

- **Domain allow lists:** You can specify exactly which web domains the agent is allowed to access.

- **Terminal approvals:** By default, the agent will always ask for your permission before executing terminal commands.
- **Sandbox mode:** The agent's file access is sandboxed to your specific project workspace. For example, it cannot rogue-access outside folders, like your personal downloads folder.

We received so many incredible questions during the broadcast that we couldn't get to all of them live. Here are the answers to the top questions we missed!

8. What is the difference between Antigravity and the VS Code plugin for Gemini?

Question from @stamy & @teronintech (YouTube): *What is the difference between Antigravity and the VS Code plugin for Gemini/Code Assistant?*
Answer: While both leverage Google's models, they are fundamentally different tools. Gemini Code Assist (the VS Code Plugin) is a traditional AI coding assistant that lives inside your existing editor, helping you with contextual suggestions and chat while you type. Antigravity, on the other hand, is a standalone, agent-first IDE. It acts as a "Mission Control" where you orchestrate autonomous agents that can plan, read and modify multiple files, run tests, and use a built-in browser to act on your behalf.

9. How do I choose between Gemini 3 Flash and Gemini 3.1 Pro?

Question from @BrentPeluso & Nicola Spreafico: *When would you use Gemini 3 FLASH vs 3.1 Pro (High)? How can I properly choose the model to use based on the task?*

Answer: Choose **Gemini 3 Flash** for speed and cost-efficiency. It is optimized for responsiveness and should be your default for high-volume tasks, quick interactive edits, and real-time systems. Choose **Gemini 3.1 Pro** when intelligence and reasoning depth are more important than speed. The 3.1 Pro model — particularly with the "High" thinking level enabled — behaves like a scaled-down version of Google's Deep Think and is meant for complex architectural refactoring, scientific tasks, and long-horizon planning.

10. How can I stop Antigravity from slowing down on large context tasks?

Question from Abdoul Aphis Traore (LinkedIn): *Why does Antigravity slow down when chat context grows larger?*

Answer: Loading too many files, rules, and tools into the agent’s context window leads to “tool bloat”, which increases latency, drives up costs, and can confuse the model. To fix this, use **Skills** (Progressive Disclosure). Skills are specialized packages of knowledge that sit dormant in your workspace and are only loaded into the agent’s context when a specific request requires them.

11. Do you have tips for interacting with agent management effectively?

Question from @jlaub (YouTube): *Do you have tips for improving how I interact with agent management and flexibility?*

Answer: The most effective workflow is to spend about 80% of your time collaborating with the agent to iterate on the “implementation plan” *before* letting the agent generate any code. If the high-level plan points to the correct files and uses the right architecture, you have a much higher likelihood of success.

12. Are there fine-grained controls for sub-agents?

Question from Mike Gugel (LinkedIn): *Are there any plans to introduce fine-grained sub-agents control?*

Answer: Antigravity is already built on a multi-agent orchestration architecture. For example, when an agent needs to use the web, the main agent delegates the exact instructions to a specialized “browser subagent” (powered by a computer-use model) whose only job is to actuate the browser and report back. You can also use the Agent Manager to manually spawn multiple independent agents working on different workspaces simultaneously.

13. How do teams collaborate and see context?

Question from Daniel Förberg & Atul Singh (LinkedIn): *How does team coordination of agents work? How can teams see each other’s context window?*

Answer: Antigravity currently doesn’t have built-in team features for shared context windows. To collaborate, teams should rely on Git or other source control systems to manage and share the artifacts produced by Antigravity. To share context seamlessly, you can ask the agent to generate a Markdown report that includes the key decisions, code snippets, and summaries of your discussion. Additionally, you can package the generated code into a zip file, or summarize multiple conversations and compile them into a single “handover” document for your teammates.

14. Can I use an ultimate ‘YOLO’ mode?

Question from @redytor (YouTube): *How about adding an ultimate ‘YOLO’ mode that runs in a sandbox without asking for permissions?*

Answer: You can already check Enable Terminal Sandbox in Agent Settings to enable terminal commands to run with sandbox restrictions on local macOS. You also have the option of “Always Proceed” (except those in the Deny list) with Terminal Command Auto Execution in Agent Settings.

15. What are the best practices for security and privacy-first workflows?

Question from @Bright2684 (YouTube): *Do you have any specific workflows or advice for security and privacy-first production-grade development?*

Answer: For maximum security, enable “Strict Mode” in Agent Settings. Set the Terminal Command Auto Execution policy to “Request Review”, configure both an **Allow List** and **Deny List Terminal Commands**, and configure a **Browser URL Allowlist** to restrict the agent's web browsing to trusted domains, helping prevent prompt injection attacks.

16. Can Antigravity run offline for privacy needs?

Question from Sandip Shrestha (LinkedIn): *Will there be a version of Antigravity that can run partially offline or with local-only RAG for high privacy needs?*

Answer: Currently, Antigravity relies heavily on cloud-based frontier models like Gemini 3.1 Pro and Gemini 3 Flash to power its complex reasoning and multi-agent orchestration. While a fully offline version isn’t available, Google does provide enterprise configurations (like VPC Service Controls) for enterprise teams that handle sensitive corporate data.

17. Can I launch an app to the App Store using Antigravity?

Question from @QuickRank-ix (YouTube): *Can I build a website and launch it with a custom domain, and build a mobile app and launch it in the Google Play or Apple store?*

Answer: Antigravity can absolutely help you design, build, and deploy the code for your applications. For instance, you can use an agent to build and deploy a serverless application directly to Google Cloud. However, the administrative processes — like purchasing a custom domain or formally submitting an application package to the Apple App Store or Google Play Store — are external actions that you must complete manually.

18. What should college graduates prioritize?

Question from Praneet V. G. (LinkedIn): *What should college graduates prioritize or put their efforts on?*

Answer: Graduates should prioritize embracing AI tools and learning to orchestrate them. AI is significantly raising the baseline for what is considered an effective engineer. Because the nature of software engineering is shifting from typing code to supervising multiple agents simultaneously, learning to break down complex problems and reviewing agent-generated plans will help you compete at a much higher level.

Want to learn more? Checkout the basics of Antigravity [video](#) and getting started [codelab](#).



Last published 1 day ago

A collection of technical articles and blogs published or curated by Google Cloud Developer Advocates. The views expressed are those of the authors and don't necessarily reflect those of Google.